

ridge-1

April 4, 2024

0.1 Predicting House Sale Prices

```
[1]: #importing necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error, r2_score
```

```
[2]: # Loading the dataset
data = pd.read_csv(r'C:\Users\ntpc\Desktop\HousePrices.csv')
data.head()
```

```
[2]:
```

| | Id | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | \ |
|---|----|------------|----------|-------------|---------|--------|-------|----------|---|
| 0 | 1 | 60 | RL | 65.0 | 8450 | Pave | NaN | Reg | |
| 1 | 2 | 20 | RL | 80.0 | 9600 | Pave | NaN | Reg | |
| 2 | 3 | 60 | RL | 68.0 | 11250 | Pave | NaN | IR1 | |
| 3 | 4 | 70 | RL | 60.0 | 9550 | Pave | NaN | IR1 | |
| 4 | 5 | 60 | RL | 84.0 | 14260 | Pave | NaN | IR1 | |

| | LandContour | Utilities | ... | PoolArea | PoolQC | Fence | MiscFeature | MiscVal | MoSold | \ |
|---|-------------|-----------|-----|----------|--------|-------|-------------|---------|--------|---|
| 0 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 | 2 | |
| 1 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 | 5 | |
| 2 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 | 9 | |
| 3 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 | 2 | |
| 4 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 | 12 | |

| | YrSold | SaleType | SaleCondition | SalePrice |
|---|--------|----------|---------------|-----------|
| 0 | 2008 | WD | Normal | 208500 |
| 1 | 2007 | WD | Normal | 181500 |
| 2 | 2008 | WD | Normal | 223500 |
| 3 | 2006 | WD | Abnorml | 140000 |
| 4 | 2008 | WD | Normal | 250000 |

[5 rows x 81 columns]

```
[3]: # getting the null values
data.isnull().sum()
```

```
[3]: Id                0
     MSSubClass        0
     MSZoning          0
     LotFrontage      259
     LotArea           0

     ...
     MoSold            0
     YrSold            0
     SaleType          0
     SaleCondition     0
     SalePrice         0
     Length: 81, dtype: int64
```

```
[4]: #checking how much percent of that column is missing
data.isnull().mean().sort_values(ascending = False)
```

```
[4]: PoolQC           0.995205
     MiscFeature      0.963014
     Alley            0.937671
     Fence            0.807534
     MasVnrType       0.597260

     ...
     ExterQual        0.000000
     Exterior2nd      0.000000
     Exterior1st      0.000000
     RoofMatl         0.000000
     SalePrice        0.000000
     Length: 81, dtype: float64
```

```
[5]: # Converting categorical variables to numerical using one-hot encoding
data = pd.get_dummies(data, columns=['MSZoning', 'Street', 'Alley', 'LotShape', '
↳ 'LandContour', 'Utilities', 'SaleType', 'SaleCondition'])

# Dropping non-numeric columns with non-numeric data
data = data.select_dtypes(include=['number'])

# Computing the correlation matrix
corr_matrix = data.corr()

# Sorting the correlation values of the target variable 'SalePrice' with other
↳ variables
corr_ser = corr_matrix['SalePrice'].sort_values(ascending=False)

print(corr_ser)
```

```
SalePrice          1.000000
OverallQual        0.790982
```

| | |
|---------------|-----------|
| GrLivArea | 0.708624 |
| GarageCars | 0.640409 |
| GarageArea | 0.623431 |
| TotalBsmtSF | 0.613581 |
| 1stFlrSF | 0.605852 |
| FullBath | 0.560664 |
| TotRmsAbvGrd | 0.533723 |
| YearBuilt | 0.522897 |
| YearRemodAdd | 0.507101 |
| GarageYrBltd | 0.486362 |
| MasVnrArea | 0.477493 |
| Fireplaces | 0.466929 |
| BsmtFinSF1 | 0.386420 |
| LotFrontage | 0.351799 |
| WoodDeckSF | 0.324413 |
| 2ndFlrSF | 0.319334 |
| OpenPorchSF | 0.315856 |
| HalfBath | 0.284108 |
| LotArea | 0.263843 |
| BsmtFullBath | 0.227122 |
| BsmtUnfSF | 0.214479 |
| BedroomAbvGr | 0.168213 |
| ScreenPorch | 0.111447 |
| PoolArea | 0.092404 |
| MoSold | 0.046432 |
| 3SsnPorch | 0.044584 |
| BsmtFinSF2 | -0.011378 |
| BsmtHalfBath | -0.016844 |
| MiscVal | -0.021190 |
| Id | -0.021917 |
| LowQualFinSF | -0.025606 |
| YrSold | -0.028923 |
| OverallCond | -0.077856 |
| MSSubClass | -0.084284 |
| EnclosedPorch | -0.128578 |
| KitchenAbvGr | -0.135907 |

Name: SalePrice, dtype: float64

```
[6]: #selecting top 10 predictors
columns = corr_ser.index[:10]
columns
```

```
[6]: Index(['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars', 'GarageArea',
        'TotalBsmtSF', '1stFlrSF', 'FullBath', 'TotRmsAbvGrd', 'YearBuilt'],
        dtype='object')
```

```
[7]: df2 = data.loc[:,columns]
```

```
[8]: df2.head()
```

```
[8]:   SalePrice  OverallQual  GrLivArea  GarageCars  GarageArea  TotalBsmtSF  \
0    208500           7       1710           2         548           856
1    181500           6       1262           2         460          1262
2    223500           7       1786           2         608           920
3    140000           7       1717           3         642           756
4    250000           8       2198           3         836          1145

   1stFlrSF  FullBath  TotRmsAbvGrd  YearBuilt
0        856         2             8        2003
1       1262         2             6        1976
2        920         2             6        2001
3        961         1             7        1915
4       1145         2             9        2000
```

```
[9]: df2.isna().sum()
```

```
[9]: SalePrice      0
OverallQual      0
GrLivArea        0
GarageCars       0
GarageArea       0
TotalBsmtSF      0
1stFlrSF         0
FullBath         0
TotRmsAbvGrd     0
YearBuilt        0
dtype: int64
```

```
[10]: #seperating X and y
X = df2.iloc[:,1:].values
y = df2.iloc[:,0].values
```

```
[11]: #splitting train and test values
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size= 0.
↪2,random_state = 0)
```

0.2 Model building

```
[12]: ridge = Ridge()

# Defining hyperparameters for grid search
param_grid = {'alpha': [0.1, 1, 10, 100]}
```

```
[13]: # Performing grid search with cross-validation
ridge_cv = GridSearchCV(ridge, param_grid, cv=5)
ridge_cv.fit(X_train, y_train)
```

```
[13]: GridSearchCV(cv=5, estimator=Ridge(), param_grid={'alpha': [0.1, 1, 10, 100]})
```

```
[14]: # Getting the best hyperparameters from grid search
best_alpha = ridge_cv.best_params_['alpha']
```

```
[15]: # Initializing Lasso Regression model with the best hyperparameters
ridge_model = Ridge(alpha=best_alpha)
ridge_model.fit(X_train, y_train)
```

```
[15]: Ridge(alpha=10)
```

```
[16]: # Making predictions on the testing data
y_pred = ridge_model.predict(X_test)
```

```
[17]: # Evaluating the model
mse = mean_squared_error(y_test, y_pred)
rmse = mse**0.5
r2 = r2_score(y_test, y_pred)
```

```
[18]: print(f'Best Alpha: {best_alpha}')
print(f'Root Mean Squared Error (RMSE): {rmse}')
print(f'R-squared (R^2): {r2}')
```

Best Alpha: 10

Root Mean Squared Error (RMSE): 50094.777084363086

R-squared (R^2): 0.6366143625748291

```
[ ]:
```